


The V8 JavaScript Engine

Design, Implementation, Testing and Benchmarking

Mads Ager, Software Engineer



Agenda

- Part 1: What is JavaScript?
- Part 2: V8 internals
- Part 3: V8 testing and benchmarking

What is JavaScript?

- The main programming language of the web
- Dynamically typed scripting language
- C-like syntax

JavaScript types

- Basic types
 - Undefined (undefined)
 - Null (null)
 - Boolean (true, false)
 - String ("a", 'a', ...)
 - Number (1, 1.23, ...)
- Object
 - Everything else

JavaScript objects

- An object is a collection of properties
- A property has a name and a value
- Properties can be added and removed on the fly

```
var o = {}; // Empty object.o.x = 42; // Adding property.o.x; // Accessing property (result: 42).o.y; // Accessing non-existing property // (result: undefined).delete o.x; // Removing property.o.x; // Accessing non-existing property.
```

JavaScript objects

- No classes in JavaScript
- All JavaScript objects are constructed from functions

```
function Point(x, y) {  
  this.x = x;  
  this.y = y;  
}
```

```
var p = new Point(1, 2);  
var q = new Point(2, 3);
```

```
var o = {}; // Basically shorthand for "new Object()"  
var a = []; // Basically shorthand for "new Array()"
```



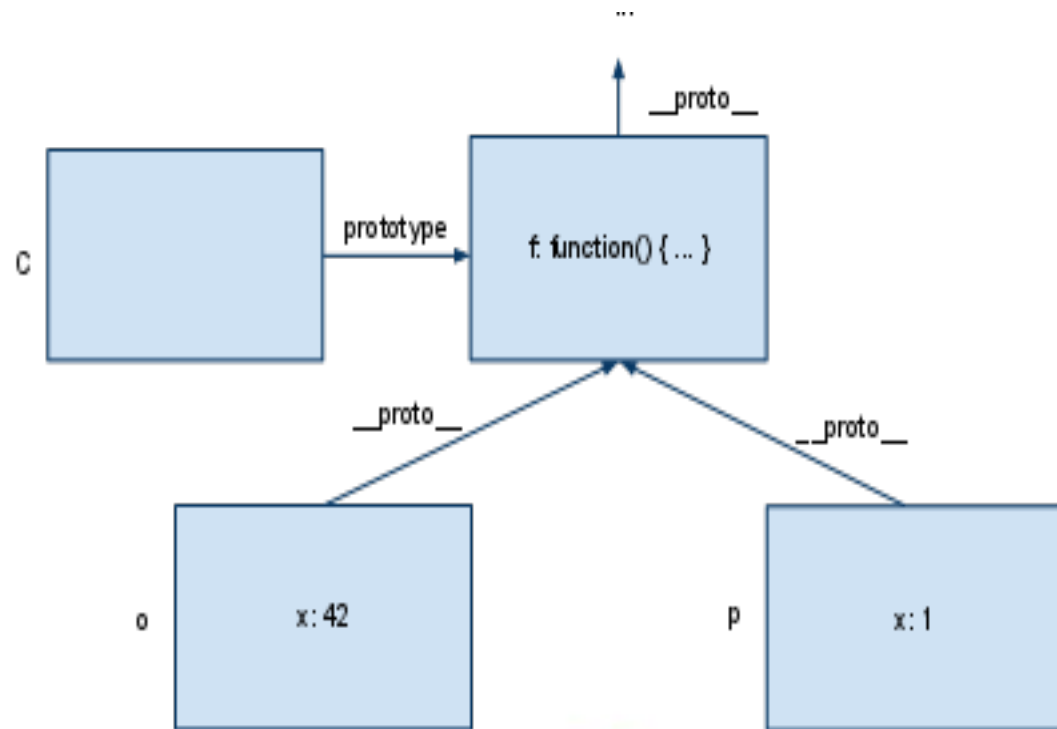
JavaScript objects

- All objects have a prototype property which can be either another object or null
- When accessing a property on an object the entire prototype chain is searched

```
function C(x) { this.x = x; }  
C.prototype.f = function() {  
  return this.x;  
}
```

```
var o = new C(42);  
o.f(); // Result: 42
```

```
var p = new C(1);  
p.f(); // Result: 1
```



JavaScript objects

- Properties that should be available on all objects can be added to `Object.prototype`
- Properties that should be available on all function objects can be added to `Function.prototype`

JavaScript objects

- Prototype chains can be used to model inheritance

```
Function.prototype.inherit = function(super) {  
  var tmpCtr = function() {};  
  tmpCtr.prototype = super.prototype;  
  this.prototype = new tmpCtr();  
  this.super_ = super;  
}  
function X(x) { this.x = x; }  
X.prototype.get_x = function() { return this.x; };  
function Y(x, y) {  
  Y.super_.call(this, x);  
  this.y = y;  
}  
Y.inherit(X);  
Y.prototype.get_y = function() { return this.y; };  
var y = new Y(1, 2);  
y.get_x(); y.get_y();
```



JavaScript types (continued)

- No types specified in the source code
- Duck typing: "If it looks like a duck and quacks like a duck, it is a duck to me"
 - All that matters is the properties, not the type

JavaScript types (continued)

- Type conversion happen at runtime to make operations make sense
- Type conversion can even be controlled by the programmer

```
var x = 1 + 2;  
x = 1 + "2";  
x = 1 + {};
```

```
var o = {};  
o.toString = function() { return "point"; }  
x = 1 + o;
```



JavaScript in browsers

- Document Object Model (DOM) objects exposed as javascript objects
- Event model: onload, onclick, ...

```
<html><body>
```

```
<script>
```

```
function changeColor(color) {  
  document.getElementById('p').style.color = color;  
}
```

```
</script>
```

```
<p id="p">
```

Changing the color of some text.

```
</p>
```

```
<button type="button" onclick="changeColor('blue')">blue</button>
```

```
<button type="button" onclick="changeColor('red')">red</button>
```

```
</body></html>
```



Agenda

- Part 1: What is JavaScript?
- Part 2: V8 internals
- Part 3: V8 testing and benchmarking

What is V8?

- JavaScript engine built from scratch in Aarhus, Denmark
- Goal was to raise the performance bar for JavaScript
- Fully open source to allow other browser vendors to use the ideas from V8





Design goals

- Make real object-oriented programs run fast
- Fast property access
- Fast function calls
- Fast and scalable memory management

The challenge

- JavaScript is very dynamic
- Properties are added and removed on the fly
- Prototype chains can be altered on the fly

Design decisions

- Introduce a notion of type called 'hidden classes' or 'maps' behind the scenes
- Generate native code
- Use a technique called inline caching for property access and function calls
- Use precise generational garbage collection

V8 memory model

- 32-bit tagged pointers
- Objects are 4-byte aligned, so two bits available for tagging
- Small 31-bit signed integers are immediate values distinguished from pointers by tags

Small integer

XXXX...XXXX0

Pointer

XXXX...XXX01

- Base JavaScript objects consists of three words

Hidden Class Pointer

Properties Pointer

Elements Pointer

Hidden classes

- Wanted to take advantage of optimization techniques from statically typed object oriented languages
- Introduced the concept of hidden classes to get there
- Hidden classes group objects that have the same structure

Hidden Classes by Example

- JavaScript objects constructed in the same way should get the same hidden class

```
function Point(x, y) {  
  this.x = x;  
  this.y = y;  
}  
var p1 = new Point(0,1);  
var p2 = new Point(2,3);
```

Hidden Classes by Example

```
function Point(x, y) {  
  this.x = x;  
  this.y = y;  
}  
var p1 = new Point(0,1);  
var p2 = new Point(2,3);
```


Point function
initial class

Class 0




Hidden Classes by Example

```
function Point(x, y) {  
  this.x = x;  
  this.y = y;  
}  
var p1 = new Point(0,1);  
var p2 = new Point(2,3);
```



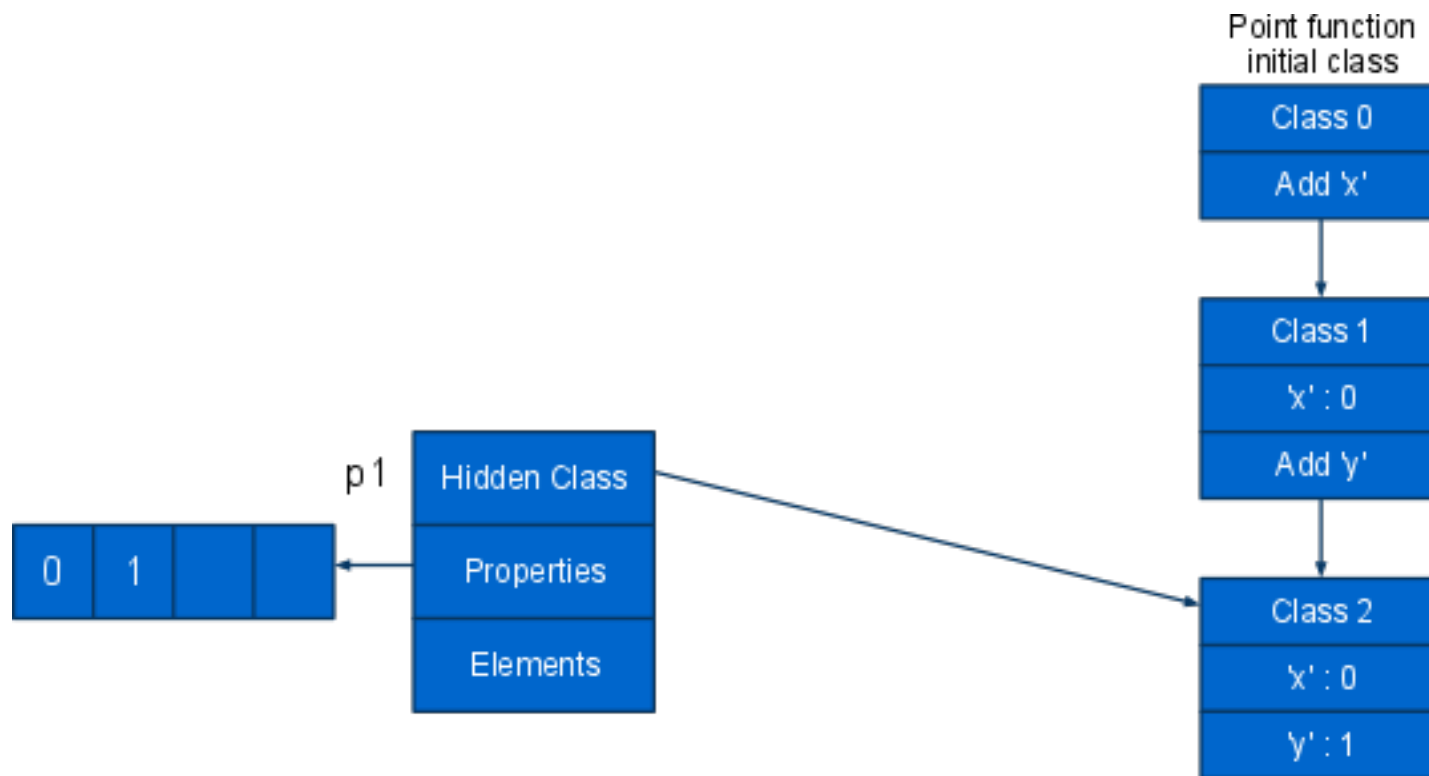
Hidden Classes by Example

```
function Point(x, y) {  
  this.x = x;  
  this.y = y;  
}  
var p1 = new Point(0,1);  
var p2 = new Point(2,3);
```



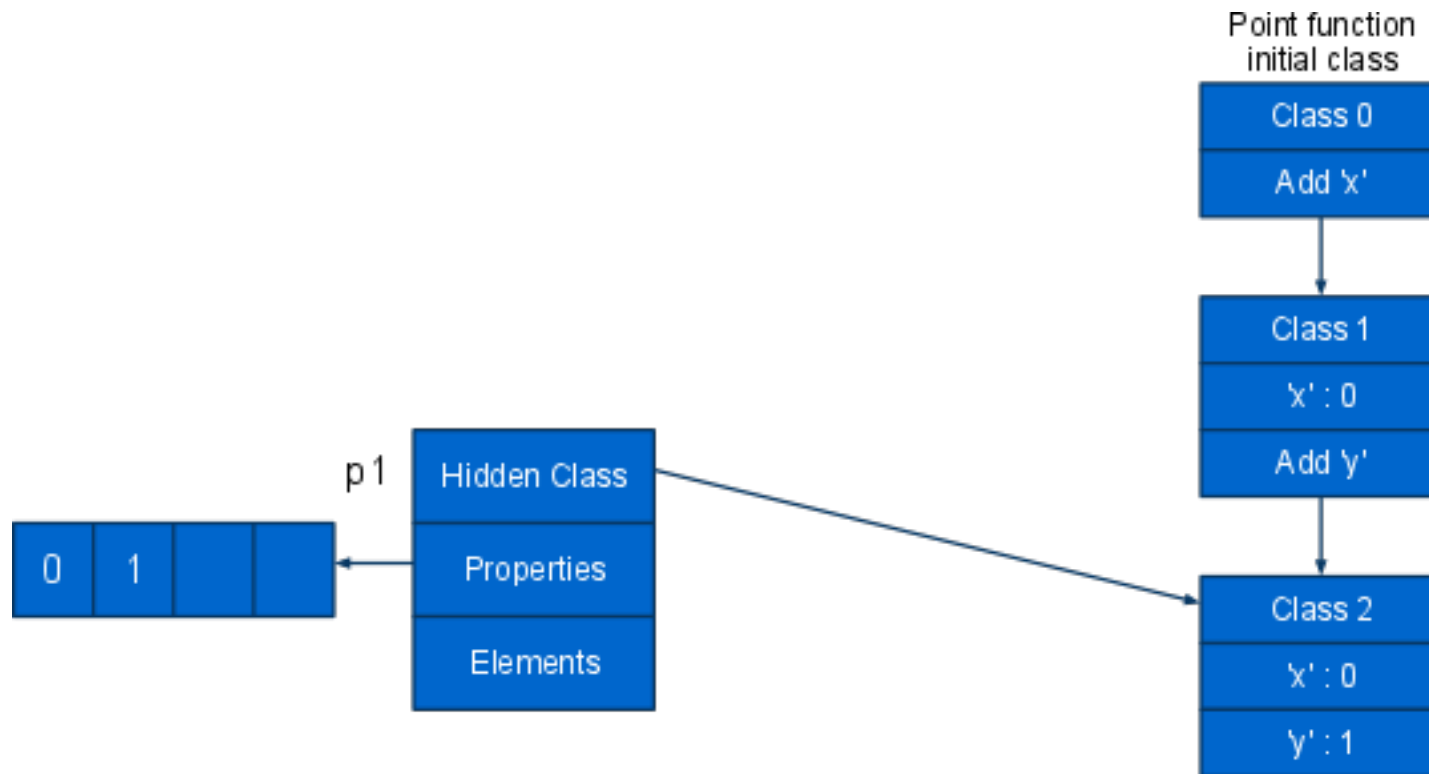
Hidden Classes by Example

```
function Point(x, y) {  
  this.x = x;  
  this.y = y;  
}  
var p1 = new Point(0,1);  
var p2 = new Point(2,3);
```



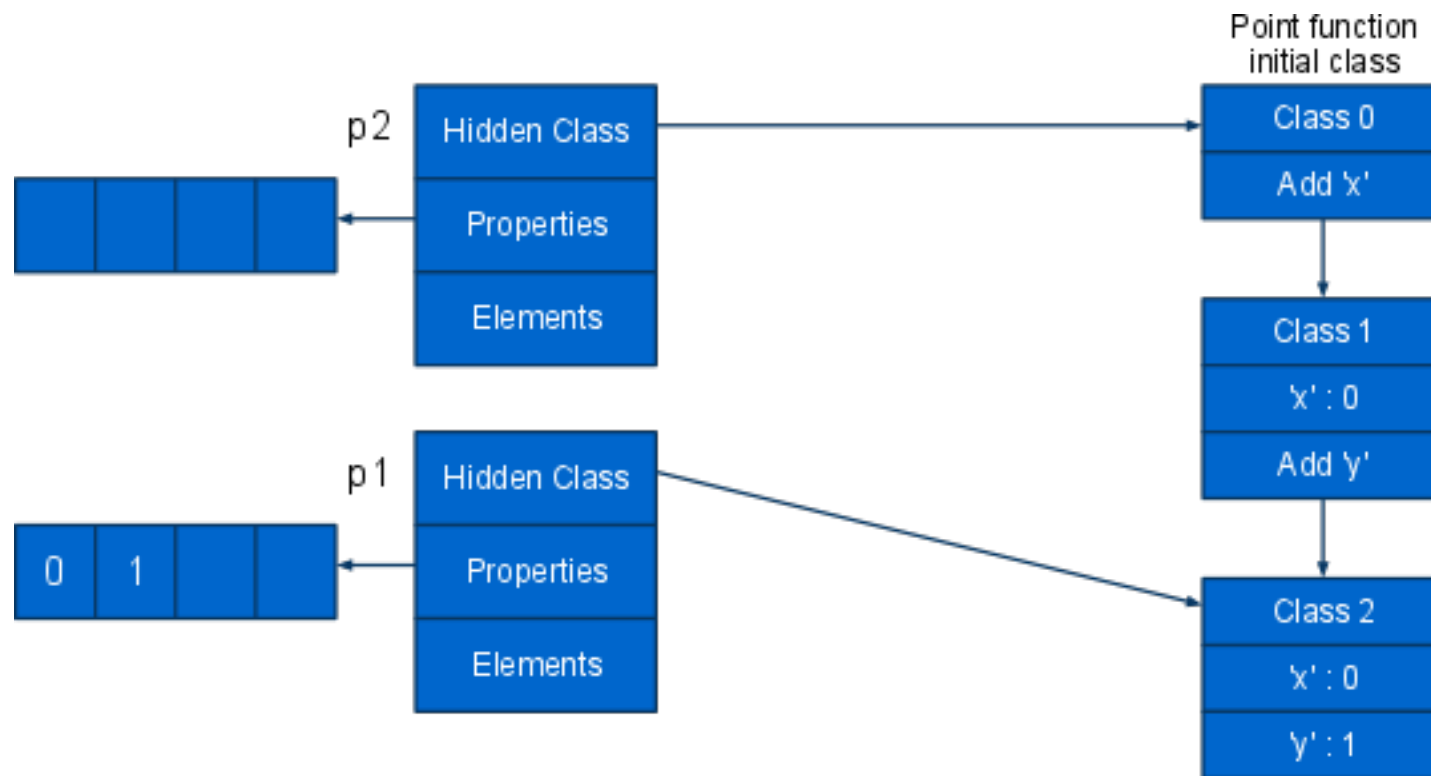
Hidden Classes by Example

```
function Point(x, y) {  
  this.x = x;  
  this.y = y;  
}  
var p1 = new Point(0,1);  
var p2 = new Point(2,3);
```



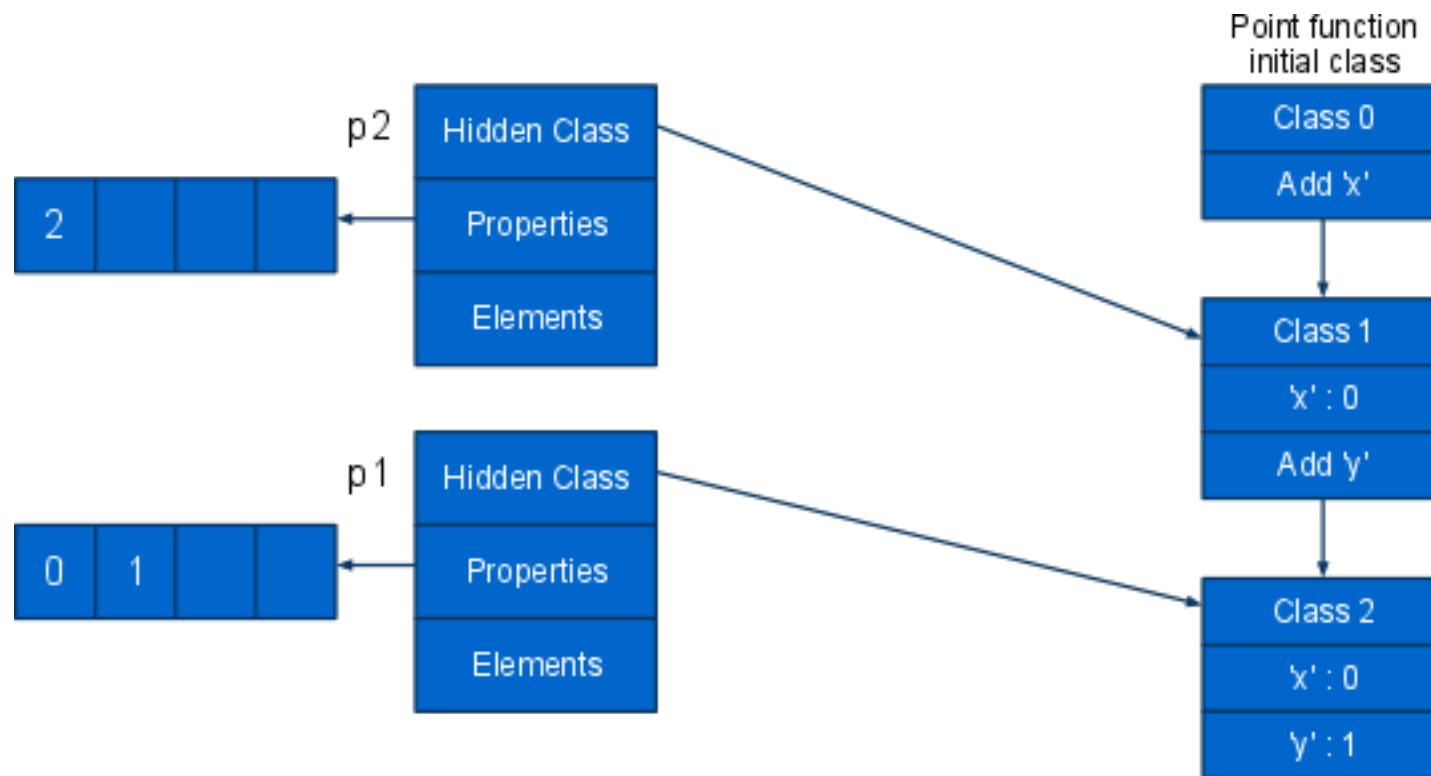
Hidden Classes by Example

```
function Point(x, y) {  
  this.x = x;  
  this.y = y;  
}  
var p1 = new Point(0,1);  
var p2 = new Point(2,3);
```



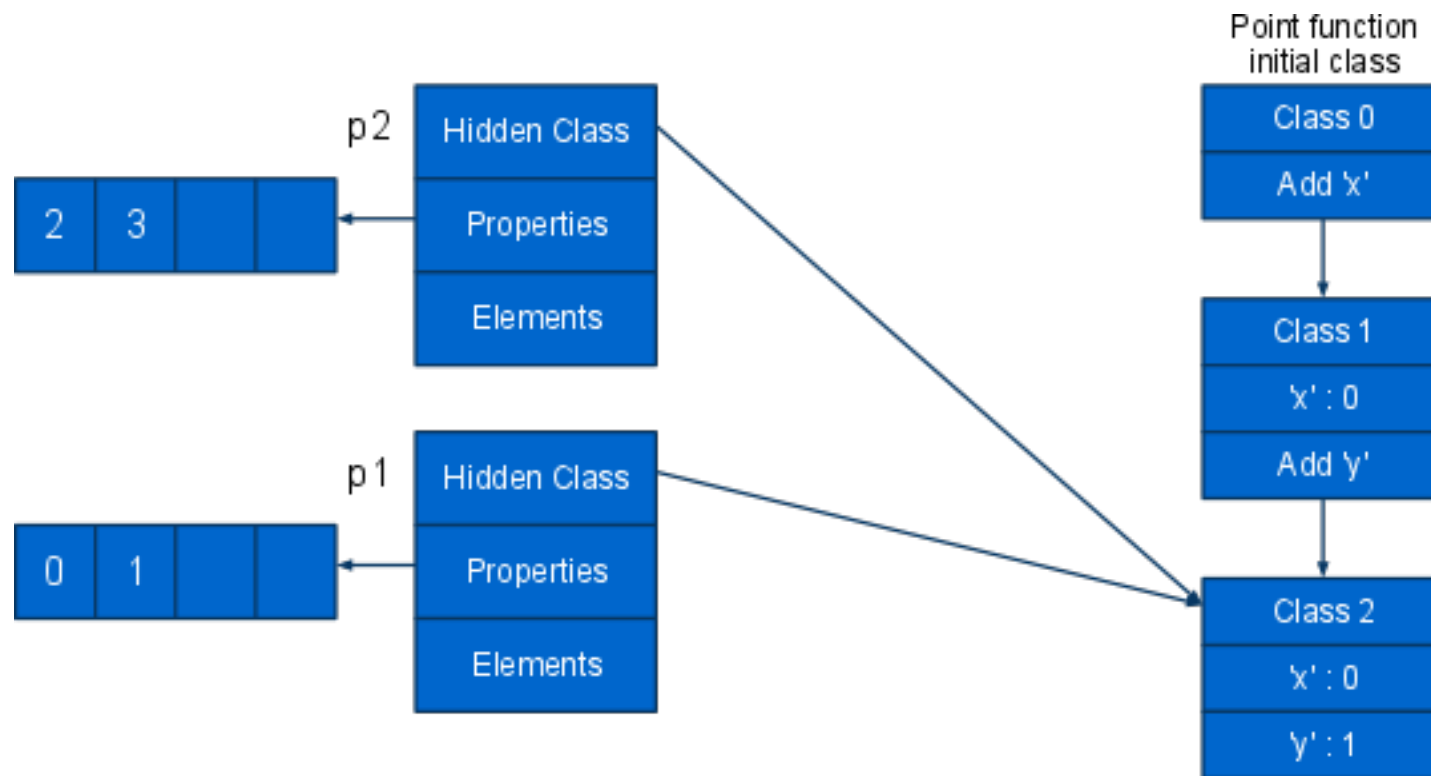
Hidden Classes by Example

```
function Point(x, y) {  
  this.x = x;  
  this.y = y;  
}  
var p1 = new Point(0,1);  
var p2 = new Point(2,3);
```



Hidden Classes by Example

```
function Point(x, y) {  
  this.x = x;  
  this.y = y;  
}  
var p1 = new Point(0,1);  
var p2 = new Point(2,3);
```



How Dynamic is JavaScript?

- In 90% of the cases, only objects having the same map are seen at an access site
- Hidden classes provides enough structure to use optimization techniques from more static object-oriented language
- We do not know the hidden class of objects at compile time
- We use runtime code generation and a technique called inline caching

Inline Caching

...



...



...

Inline Caching

...



...



...

Inline Caching

...

load 'x'



...

➔ load 'y'



...

Inline Caching

...

load 'x'



...

➔ load 'y'



...

Monomorphic Load Inline Cache

0xf7c0d32d (size = 37):

0 mov eax,[esp+0x4]

4 test al,0x1

6 jz 32

12 cmp [eax+0xff],0xf78fab81

19 jnz 32

25 mov ebx,[eax+0x3]

28 mov eax,[ebx+0x7]

31 ret

32 jmp LoadIC_Miss



Monomorphic Load Inline Cache

0xf7c0d32d (size = 37):

0 mov eax,[esp+0x4] ; receiver load

4 test al,0x1 ; object check

6 jz 32

12 cmp [eax+0xff],0xf78fab81

19 jnz 32

25 mov ebx,[eax+0x3]

28 mov eax,[ebx+0x7]

31 ret

32 jmp LoadIC_Miss

Monomorphic Load Inline Cache

0xf7c0d32d (size = 37):

0 mov eax,[esp+0x4] ; receiver load

4 test al,0x1 ; object check

6 jz 32

12 cmp [eax+0xff],0xf78fab81 ; class check

19 jnz 32

25 mov ebx,[eax+0x3]

28 mov eax,[ebx+0x7]

31 ret

32 jmp LoadIC_Miss

Monomorphic Load Inline Cache

0xf7c0d32d (size = 37):

0 mov eax,[esp+0x4] ; receiver load

4 test al,0x1 ; object check

6 jz 32

12 cmp [eax+0xff],0xf78fab81 ; class check

19 jnz 32

25 mov ebx,[eax+0x3] ; load properties

28 mov eax,[ebx+0x7] ; load property

31 ret

32 jmp LoadIC_Miss

Monomorphic Load Inline Cache

0xf7c0d32d (size = 37):

0 mov eax,[esp+0x4] ; receiver load

4 test al,0x1 ; object check

6 jz 32

12 cmp [eax+0xff],0xf78fab81 ; class check

19 jnz 32

25 mov ebx,[eax+0x3] ; load properties

28 mov eax,[ebx+0x7] ; load property

31 ret

32 jmp LoadIC_Miss ; fallback to
; generic lookup

Inline Cache States

- Three inline cache states
 - Uninitialized
 - Monomorphic
 - Megamorphic
- In the megamorphic state a cache of generated stubs is used
- Inline caches are cleared on full garbage collections
 - Allows us to get rid of unused code stubs
 - Gives all inline caches a new chance to hit the monomorphic case

Memory Management

- JavaScript is garbage collected
- Users don't have to worry about deallocating memory, the virtual machine will get rid of unused memory

Memory Management Basics

- Memory is allocated from a piece of memory called the heap which is controlled by the virtual machine
- When there is too little memory left to fulfill an allocation request a garbage collection is triggered
- Basic garbage collection
 - Find all objects reachable from a root set (statics, pointers on the stack)
 - Allow the rest of memory to be reused

Mark-Sweep

- Two phases
- Mark
 - Starting from the root set perform graph traversal marking every object that is reachable
- Sweep
 - Sweep through the entire heap from beginning to end
 - For each object encountered
 - Remove mark if marked and leave alone
 - If not marked make space available for allocation (for instance by adding to a free list)

Copying

- Two separate spaces of the same size
- Always allocate from one space by just moving a pointer
- To collect garbage everything that is reachable is copied to the other space and allocation continues there
- Pros: Fast. Only live data is touched.
- Cons: Wastes space.

Efficient Memory Management

- Precise generational garbage collection
- Two generations
 - Young generation is one small, contiguous space that is collected often
 - Old generation is divided into a number of spaces that are only occasionally collected
 - Code space (executable)
 - Map space (hidden classes)
 - Large object space (>8K)
 - Old data space (no pointers)
 - Old pointer space
- Objects are promoted from young to old generation if they survive

Types of Garbage Collection

- Scavenge
 - Copying collection of only the young generation
 - Pause times ~2ms
- Full non-compacting collection
 - Mark-Sweep collection of both generations
 - Free memory gets added to free lists
 - Might cause fragmentation
 - Pause times ~50ms
- Full compacting collection
 - Mark-Sweep-Compact collection of both generations
 - Pause times ~100ms

Agenda

- Part 1: What is JavaScript?
- Part 2: V8 internals
- Part 3: V8 testing and benchmarking

Testing



- V8 is shipped in Google Chrome every week
- We need V8 to be stable
- We need to push performance and be aggressive
- Therefore, we need a lot of testing

Testing

- For each commit to the V8 repository
 - V8 tests (>1200 tests)
 - Mozilla JS tests (>1900 tests)
 - Sputnik (>5000 tests)
 - ES5 conformance tests (>1200 tests)
- Many test configurations
 - Linux x { ia32, x64, arm simulator } x { release, debug }
 - Windows x { ia32, x64 } x { release, debug }
 - Mac x ia32 x { release, debug }

Testing

- For each V8 commit
 - WebKit layout tests run in Chromium test shell
- Push new version to Chromium twice a week
- After each push a lot of additional testing goes on on the Chromium buildbots

<http://build.chromium.org/buildbot/v8>
<http://build.chromium.org>



Testing


- Testing is crucial to making progress
- Having good test coverage allows us to make aggressive changes with less fear of breaking stuff
- When fixing a bug you add a new test
- When implementing a new feature you add a new test

Benchmarking

- We want to push JavaScript performance
- Need to track how we are doing
- We have built an infrastructure that runs a lot of benchmarks on every commit to the V8 repository
- Easy to keep track of progress and easy to spot regressions


Performance tracking

V8BenchmarkSuite-geometric-mean




Performance tracking

SunSpider-sum



Performance

- Goal was to push performance
- Performance is increased on every release of Google Chrome



Questions?

<http://v8.googlecode.com>

<http://chromium.org>



Google™